


Week 4

Process and Thread Management – Operating System Support

Assignment Project Exam Help



Reference:

Chapter 7

<https://powcoder.com>

Distributed Systems: Concepts and Design

Coulouris, Dollimore, Kindberg and Blair

Edition 5, © Addison Wesley 2011

Add WeChat powcoder

Learning Objectives

- Explain what a modern operating system does in support of distributed applications and middleware.
 - Network operating systems
 - Distributed operating systems
 - Supporting distributed applications by the combination of middleware and NOSs
- Recognise OS abstractions for resource management.
 - Process execution environment
 - Multiple processes and threads mechanisms.

Learning Objectives

- Develop Java multi-thread programming for client/server applications.
- Appraise interprocess communication and invocation in terms of:
 - Performance issues and factors affecting performance
 - Potential OS or middleware supports for performance
 - Synchronous and asynchronous invocations

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

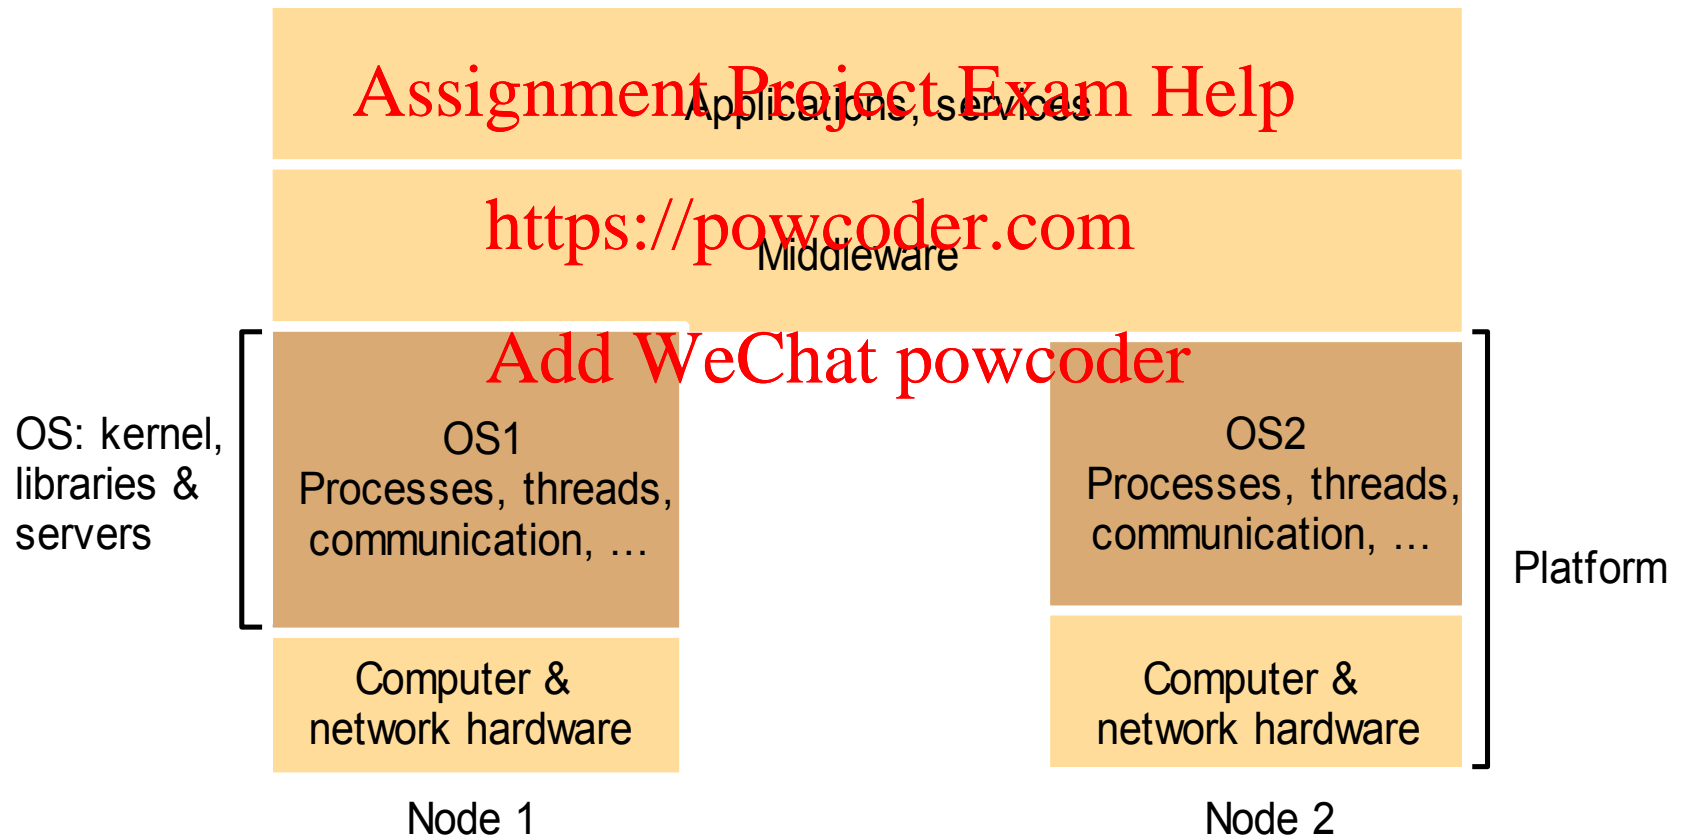
Middleware and NOSs

□ The Middleware Layer

- The client application invokes operations on another process, which are often on another node.
- Middleware provides remote invocation between objects or processes at the nodes of a distributed system.
- The requirements of middleware are to be met by the operating systems.
 - Efficient and robust access to physical resources.
 - Flexibility to implement a variety of resources-management policies.

Middleware and NOSs

□ The System layers



Middleware and NOSs

□ The operating systems

- An operating system provides abstractions of the underlying physical resources (processors, memory, communication, and storage media).
- It simplifies, protects and optimizes the use of resources.
- The *network operating systems* have a networking capability built into them and so can be used to access remote resources.
- Examples of NOSs include UNIX, Windows, and Linux.

Middleware and NOSs

□ The operating systems

- Access is network-transparent for some types of resources, such as files.
- Nodes running a network operating system retain autonomy in managing their own processing resources.
- With a network operating system, there are multiple system images/views, one per node.
- An operating system that produces a single system image/view for all resources on all nodes is called a *distributed operating system*.

Middleware and NOSs

□ The operating systems

- A DOS presents users (and applications) with an integrated computing platform that hides the individual computers.
- A DOS has control over all of the nodes (computers) in the network and allocates their resources to tasks without users' involvement.
- With a DOS, the user doesn't know (or care) where his programs are running.
- Examples include:
 - Cluster computer systems
 - V system, Sprite, Globe OS

Middleware and NOSs

□ The operating systems

□ Two reasons hinder the popularity of DOSs

- Users are not willing to adopt a new operating system that will not run their legacy applications.
- Users are not willing to lose autonomy by adopting DOSs.

□ The combination of middleware and NOSs is a solution of autonomy and network-transparency.

□ The Middleware

- Runs on a variety of OS-hardware combinations at nodes of a distributed system.
- Utilises local resources to implement its mechanism for remote invocations between objects or processes at the nodes.

□ Users will only be satisfied if their middleware-OS combination has good performance.

Middleware and NOSs

□ The operating systems

□ A NOS provides the following interfaces:

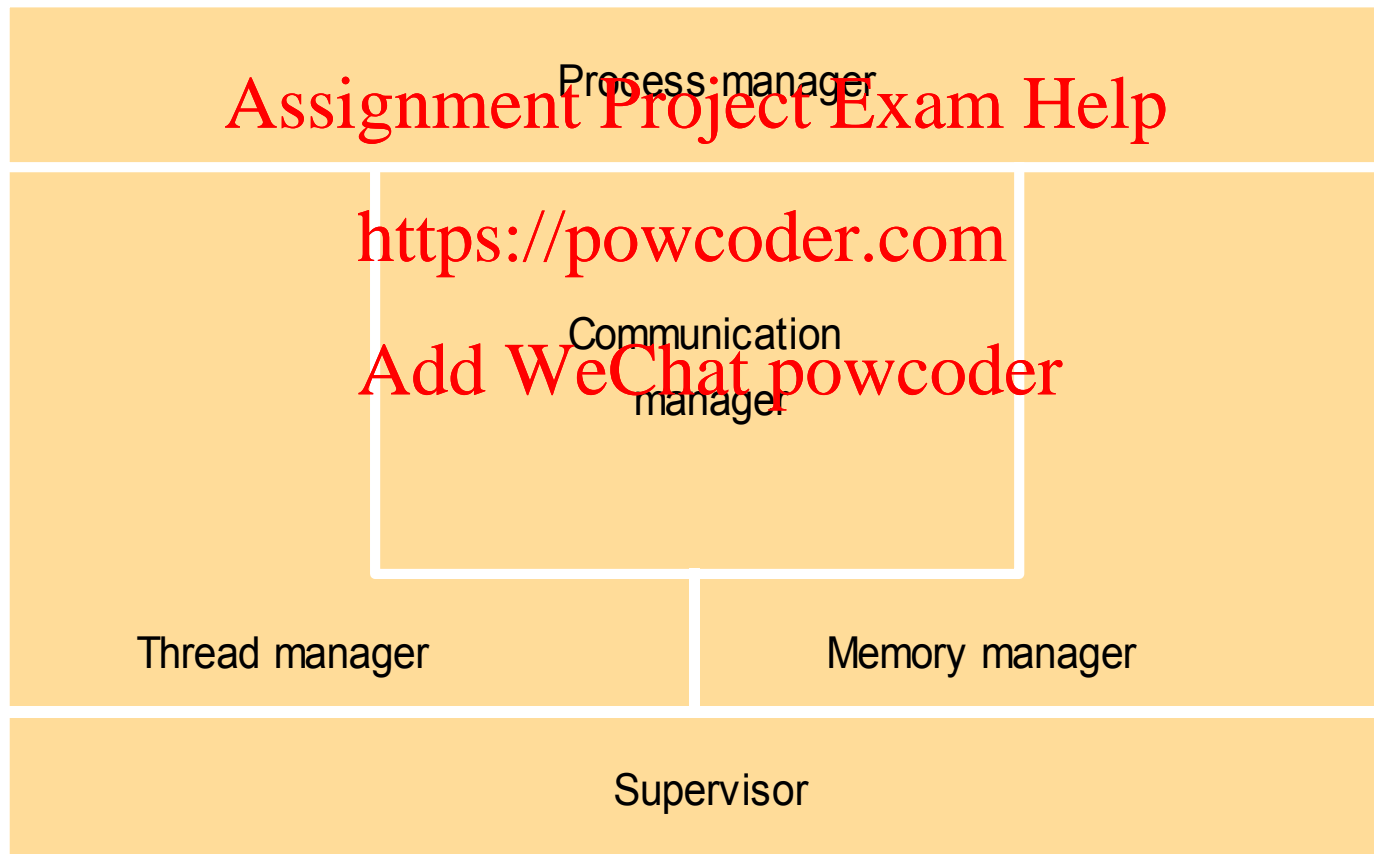
- Encapsulation of the basic resources
- Protection of the resources used by applications
- Concurrent processing to enable applications to complete their work concurrently

□ A NOS uses a combination of libraries, kernels and servers to perform

- Communication
- Scheduling
- Process management
- Thread management
- Memory management

Middleware and NOSs

□ The OS components



Resource Protection

- Resources require protection from illegitimate accesses.
 - Maliciously contrived codes (security system cares)
 - Benign codes with bugs (OS cares)
 - Examples like file read and write attributes
- To protect resources
 - The OS kernel runs in supervisor (privileged) mode and has complete privileges for physical resources.
 - The kernel arranges other process in user (unprivileged) mode for execution.

Resource Protection

□ To protect resources

- The kernel sets up address spaces to protect itself and other processes.
- An address space is a collection of ranges of virtual memory locations.
- Access rights are applied to address space.
- A process can not access memory outside its address space.
- A process can safely transfer to the kernel address space via a system call trap to execute the kernel codes.

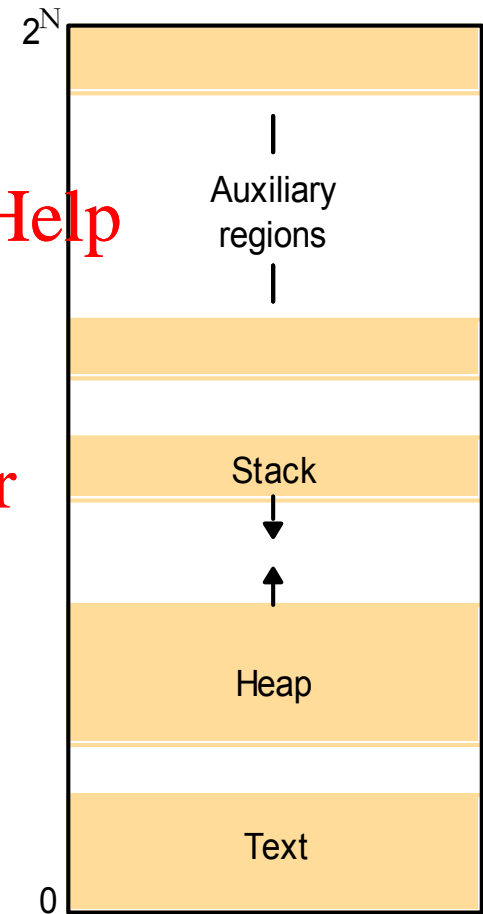
Process and Thread

- A process consists of an execution environment together with one or more threads.
- An execution environment is a unit of resource management, consisting of
 - An address space
 - Thread synchronisation and communication resources (semaphores and sockets)
 - High-level resources (open files)

Process and Thread

□ A address space consists of the following regions.

- The text region contains the program codes.
- A heap is initialised by values in the program binary file.
- A stack contains dynamically created values.



Process and Thread

- In distributed systems, the creation of a new process can be on a local node or on a remote node to utilise remote resources.
- The transfer policy determines whether to situate a new process locally or remotely.
- The location policy determines which node should host the new process selected for transfer.
 - Static: based on expected long-term characteristics.
 - Adaptive: considering the current system state.

Process and Thread


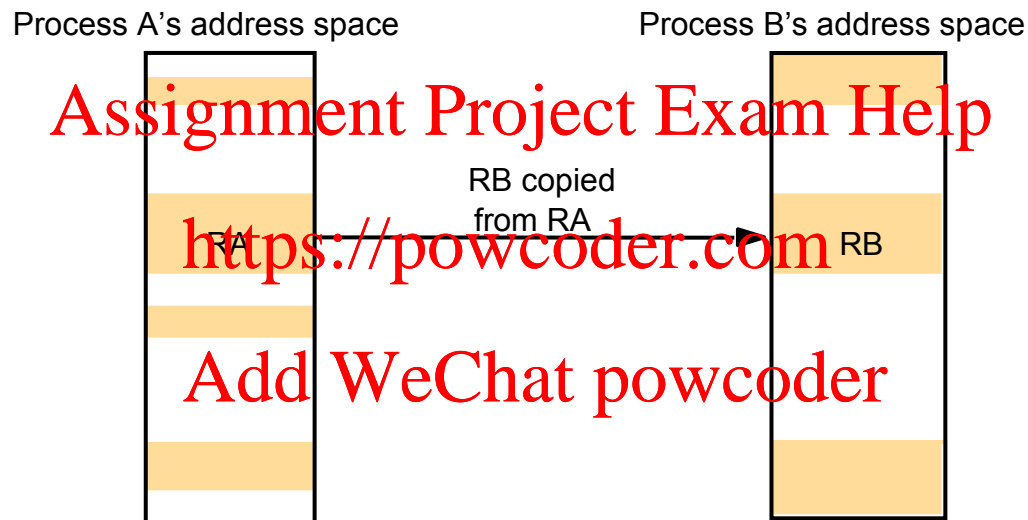
- Load-sharing systems can be:
 - Centralised
 - one load manager
 - Hierarchical
 - Managers make process allocation decisions as far down the trees as possible
 - Managers may transfer processes to one another via a common ancestor
 - Decentralised
 - Nodes exchange information with one another directly to make allocation decision.

Process and Thread

- Load-sharing systems can be:
 - Sender-initiated
 - Receiver-initiated
- When a new process is created, it requires a execution environment.
 - Initialised from an executable file
 - Copied from an existing environment, such as UNIX fork semantics
- Migratory load-sharing systems can shift load at any time, not just when a new process is created.

Process and Thread

- Copy-on-write is an optimization for copying execution environment.

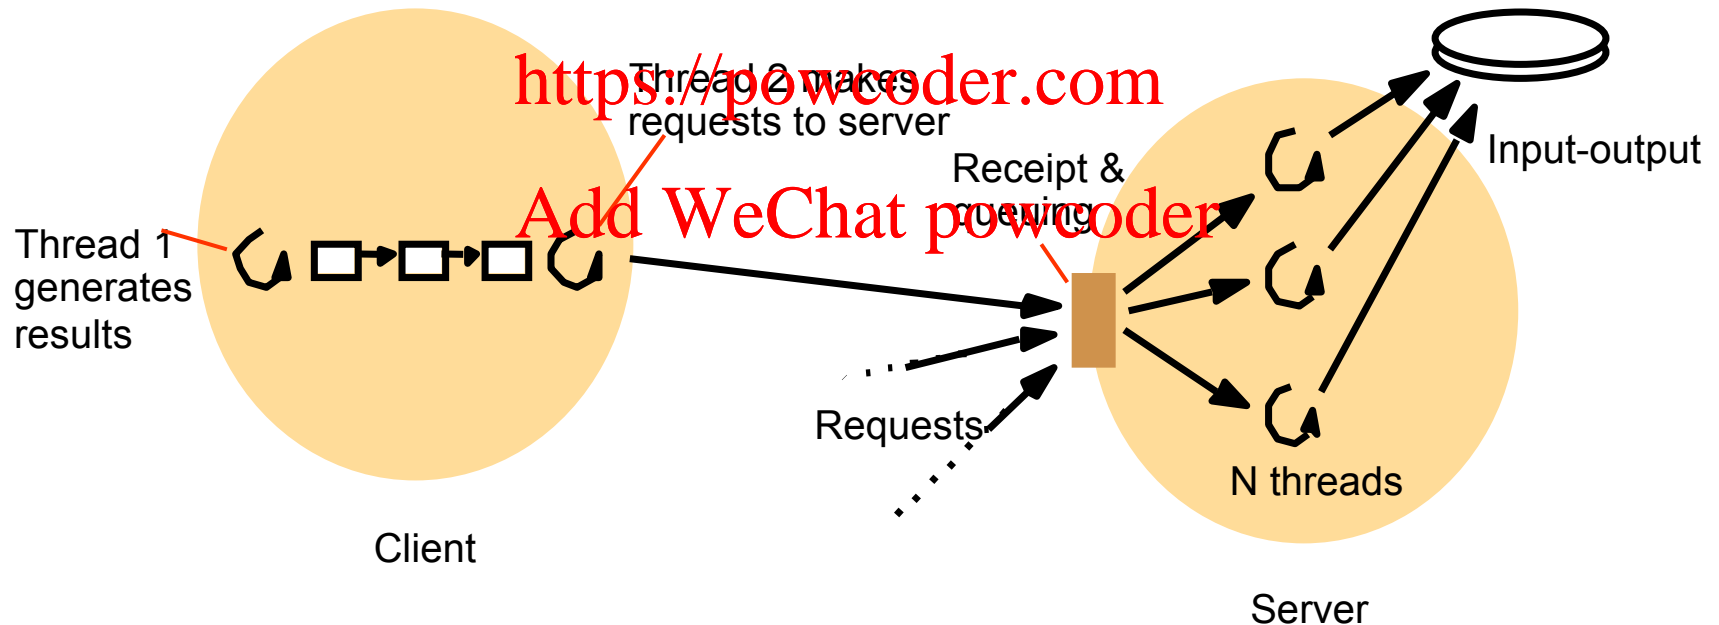


Process and Thread

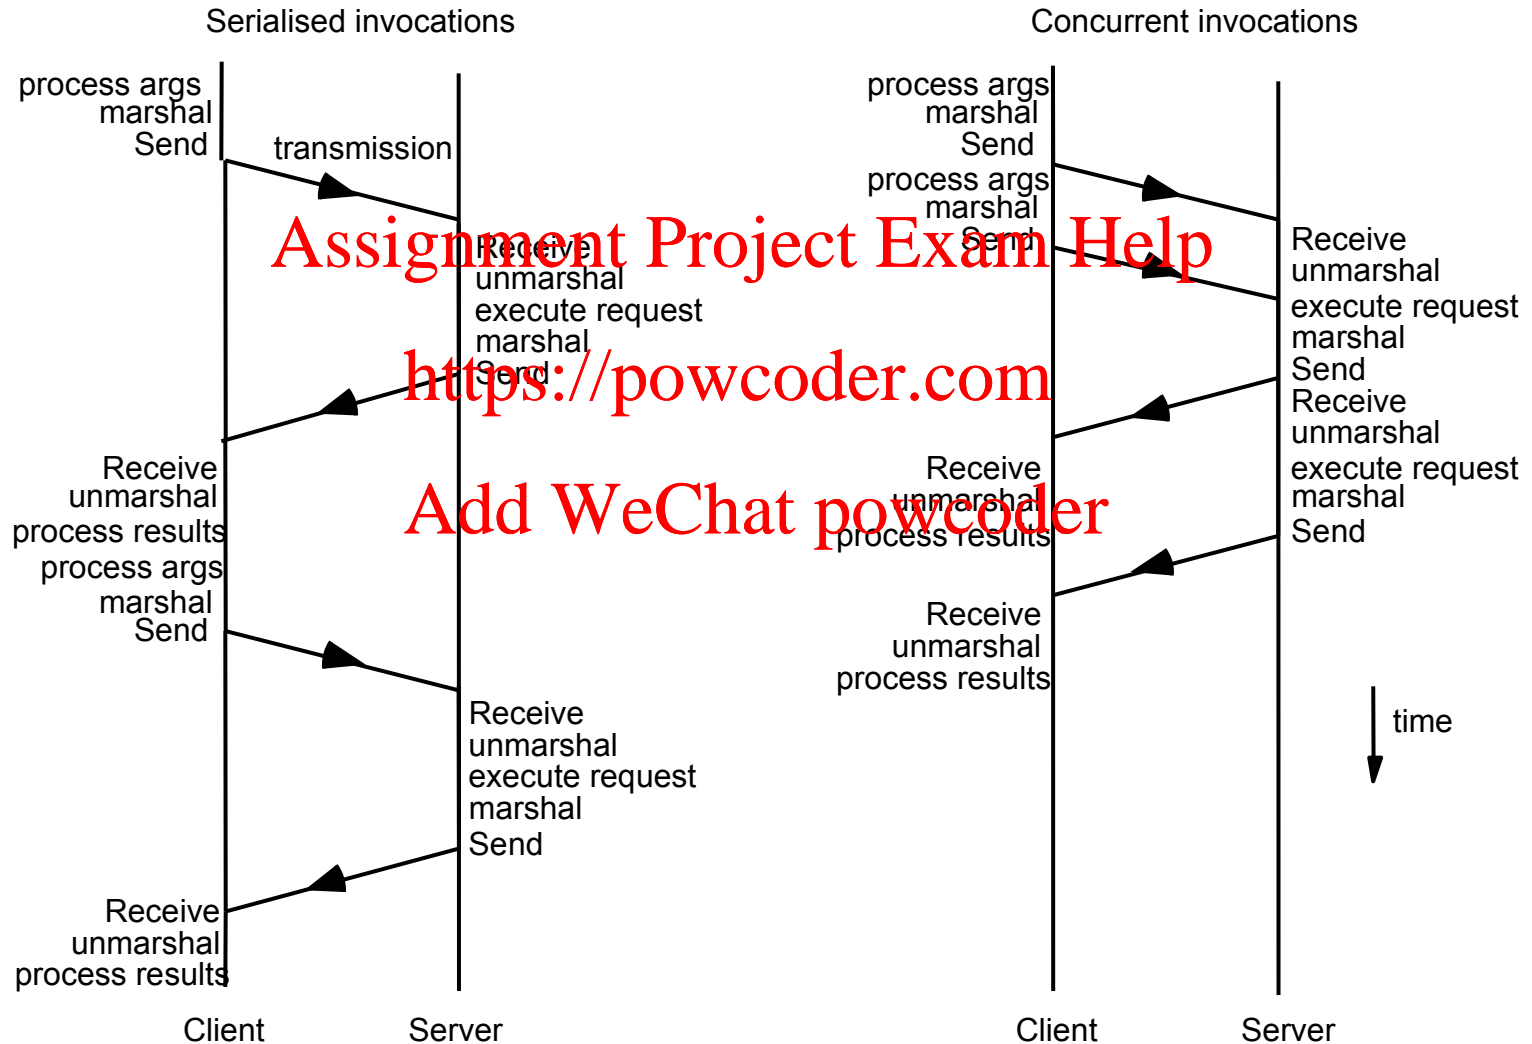
- Threads are sometimes called *lightweight processes* within a process.
- Threads share the creator process' resources, including memory and open files.
- Threads make use of concurrency to increase processing efficiency.
 - The client has two threads, one for preparing requests; the other for sending requests.
 - The server has a pool of threads, each of which removes a request from the queue and process it.

Process and Thread

□ Client and server with multiple threads



Process and Thread



Process and Thread

- The architecture for multi-threaded servers concerns the various ways of mapping requests to threads within a server.
 - The **worker pool** architecture uses a fixed size pool of worker threads.
 - The **thread-per-request** architecture creates a worker thread for each request.
 - The **thread-per-connection** architecture associates a worker thread with each connection.
 - The **thread-per-object** architecture associates a worker thread with each remote object.


Process and Thread

□ The architecture for multi-threaded servers


Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder



a. Thread-per-request



b. Thread-per-connection



c. Thread-per-object

Process and Thread

□ Threads versus multiple processes

- Creating a thread is (much) cheaper than creating a process (10-20 times).
- Switching to a different thread in same process is (much) cheaper (5-50 times)
- Threads within same process can share data and other resources more conveniently and efficiently (without copying or messaging)
- Threads within a process are not protected from each other.

Process and Thread

□ Threads within client

- The web browser is a good example of using threads within client.
- A web pages typically contains several images.
- The browser has to fetch each of these images in a separate HTTP GET request.
- The browser does not need to obtain the images in a particular sequence.
- The browser can make concurrent request by using multiple threads.
- At the same time the main thread can continue its own tasks such as image rendering.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Process and Thread

□ Threads versus multiple processes

Assignment Project Exam Help

Execution environment

Thread

Address space tables

<https://powcoder.com>

Saved processor registers

Communication interfaces, open files

Priority and execution state (such as

Add WeChat powcoder
BLOCKED)

Semaphores, other synchronization
objects

Software interrupt handling information

List of thread identifiers

Execution environment identifier

Pages of address space resident in memory; hardware cache entries

Process and Thread

- Threads programming
 - Threads programming is concurrent programming.
 - Thread synchronization is necessary to maintain data consistency.
 - A thread has its own lifetime from creation (such as using **new** operation in Java) to end (such as its **destroy()** method is called in Java).
 - Threads can be assigned with priorities, with higher priority thread has more chance to be scheduled for execution than a lower priority one.

Process and Thread

□ Threads programming

□ Threads scheduling can be

- Preemptive: a thread may be suspended at any time to make a way to another thread.
- Non-preemptive: a thread may call the threading system to yield processor.

□ Java Threads programming

- Java provides methods for creating, destroying and synchronizing threads.
- A Java thread is defined by extending the class ***Thread*** or implementing interface: ***Runnable***.
- A thread is run by calling its ***start()*** methods, which calls its ***run()*** method.
- The functional methods are defined in the ***run()*** method.

Process and Thread

□ Java threads constructor and management methods.

Thread(ThreadGroup group, Runnable target, String name)

Creates a new thread in the *SUSPENDED* state, which will belong to *group* and be identified as *name*; the thread will execute the *run()* method of *target*.

setPriority(int newPriority), getPriority()

Set and return the thread's priority.

run()

A thread executes the *run()* method of its target object, if it has one, and otherwise its own *run()* method (*Thread* implements *Runnable*).

start()

Change the state of the thread from *SUSPENDED* to *RUNNABLE*.

sleep(int millisecs)

Cause the thread to enter the *SUSPENDED* state for the specified time.

yield()

Enter the *READY* state and invoke the scheduler.

destroy()

Destroy the thread.

Process and Thread

□ Java threads synchronization calls

thread.join(int millisecs)

Blocks the calling thread for up to the specified time until *thread* has terminated.

thread.interrupt()

Interrupts *thread*: causes it to return from a blocking method call such as *sleep()*.

object.wait(long millisecs, int nanosecs)

Blocks the calling thread until a call made to *notify()* or *notifyAll()* on *object* wakes the thread, or the thread is interrupted, or the specified time has elapsed.

object.notify(), object.notifyAll()

Wakes, respectively, one or all of any threads that have called *wait()* on *object*.

Process and Thread

□ Java threads example

- This example consists of a client program and a server program, which are modified versions of the ~~TCP Client~~ and ~~TCP Server~~ on page 156-157.
<https://powcoder.com>
~~Add WeChat~~ ~~powcoder~~
- The modification is to demonstrate the *thread-per-connection* architecture
- In this example, the server dynamically creates one thread for each client connection, which accepts multiple requests and replies responses.
- In this demonstration, we use two clients, each of which sends 5 requests and receives 5 responses in a single connection.

Process and Thread

□ The TCP Client

```
import java.net.*;
import java.io.*;
public class TCPClient {
    public static void main (String args[]) {
        //arguments supply message, hostname of destination and client ID
        Socket s=null;
        try{
            int serverPort=7896;
            s=new Socket(args[1], serverPort);
            DataInputStream in = new DataInputStream(
                s.getInputStream());
            DataOutputStream out =new DataOutputStream(
                s.getOutputStream());
            for (int i=1; i<=5 ;i++ ){
                String request="Client "+args[2]+" : "+args[0]+" "+i;
                out.writeUTF(request);
                String data=in.readUTF();
                System.out.println("Received: "+ data) ;
            }
        }
```



This program continues on the next slide

Process and Thread

□ The TCP Client

```
}catch (UnknownHostException e){
    System.out.println("Sock:"+e.getMessage());
}catch (EOFException e){
    System.out.println("EOF:"+e.getMessage());
}catch (IOException e){
    System.out.println("IO:"+e.getMessage());
}finally {
    if(s!=null)
        try {s.close();}
        catch (IOException e){
            System.out.println("close:"+e.getMessage());}}
}
```

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

The end of this program

Process and Thread

□ The TCP Server

```
import java.net.*;
import java.io.*;
public class TCPServer {
    public static void main (String args[]) {
        try{
            int serverPort=7896;
            ServerSocket listenSocket=new
                ServerSocket(serverPort);

            int i=0;
            while(true) {
                Socket clientSocket=listenSocket.accept();
                Connection c = new Connection(clientSocket, i++);
                System.out.println("Thread " +i+ " is created");
            }
        } catch(IOException e) {
            System.out.println("Listen :"+e.getMessage());
        }
    }
}
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



This program continues on the next slide

Process and Thread

□ The TCP Server

```
class Connection extends Thread {
    DataInputStream in;
    DataOutputStream out;
    Socket clientSocket;
    int thrdn;
    public Connection (Socket aClientSocket, int tn){
        try {
            thrdn=tn;
            clientSocket = aClientSocket;
            in=new DataInputStream(
                clientSocket.getInputStream());
            out=new DataOutputStream(
                clientSocket.getOutputStream());
            this.start();
        } catch (IOException e) {
            System.out.println("Connection: \"
                                +e.getMessage());
        }
    }
```



This program continues on the next slide

Process and Thread

□ The TCP Server

```
public void run() {
    try {
        String data;
        while ((data=in.readUTF())!=null) {
            out.writeUTF("Reply to "+data+ " from thread "
                +Thread.currentThread().getName()+" on server");
        }
    } catch (EOFException e) {
        System.out.println("EOF:"+e.getMessage());
    } catch (IOException e) {
        System.out.println("IO:"+e.getMessage());
    } finally {
        try {clientSocket.close();}
        catch (IOException e) { /*close failed*/ }
    }
}
```

The end of this program

Process and Thread

- Output from the *TCPClient* and the *TCPServer*

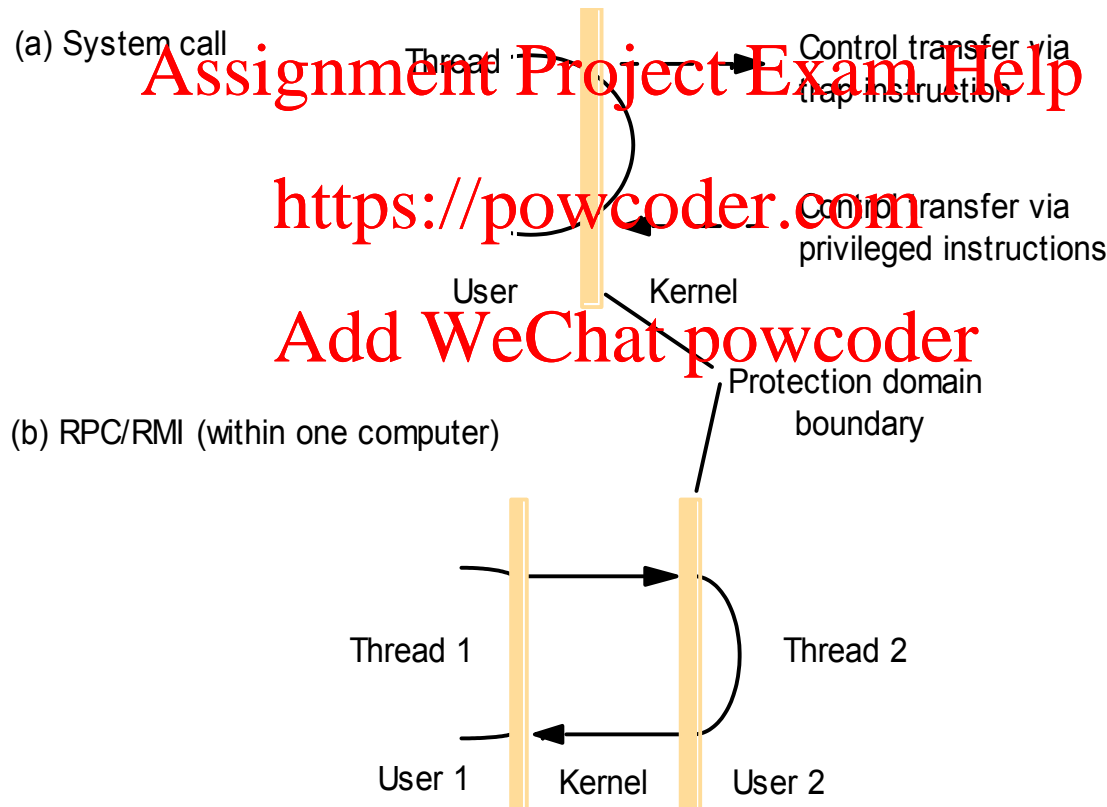
```
Administrator: Command Prompt - java TCPServer
D:\DistributedSystems\Week4\Examples>
D:\DistributedSystems\Week4\Examples>
D:\DistributedSystems\Week4\Examples>
D:\DistributedSystems\Week4\Examples>
D:\DistributedSystems\Week4\Examples>java TCPServer
Thread 1 is created
EOF:null
Thread 2 is created
EOF:null
```

```
Administrator: Command Prompt
D:\DistributedSystems\Week4\Examples>java TCPClient Hello localhost 1
Received: Reply to Client 1: Hello 1 from thread 0 on server
Received: Reply to Client 1: Hello 2 from thread 0 on server
Received: Reply to Client 1: Hello 3 from thread 0 on server
Received: Reply to Client 1: Hello 4 from thread 0 on server
Received: Reply to Client 1: Hello 5 from thread 0 on server
D:\DistributedSystems\Week4\Examples>
```

```
Administrator: Command Prompt
D:\DistributedSystems\Week4\Examples>
D:\DistributedSystems\Week4\Examples>java TCPClient Hello localhost 2
Received: Reply to Client 2: Hello 1 from thread 1 on server
Received: Reply to Client 2: Hello 2 from thread 1 on server
Received: Reply to Client 2: Hello 3 from thread 1 on server
Received: Reply to Client 2: Hello 4 from thread 1 on server
Received: Reply to Client 2: Hello 5 from thread 1 on server
D:\DistributedSystems\Week4\Examples>
```

Communication and Invocation

□ Invocations between address spaces




Communication and Invocation

□ Invocations between address spaces

Assignment Project Exam Help

(c) RPC/RMI (between computers)



Communication and Invocation

- The performance of RPC and RMI mechanisms is critical for effective distributed systems.
- A local *null procedure call* is less than 1 microseconds.
- A Remote *null procedure call* is about 10 milliseconds.
 - Network time (involving about 100 bytes transferred, at 100 megabits/sec.) accounts for only 0.01 millisecond.
 - The remaining delays must be in OS and middleware latency.

Communication and Invocation

- Factors affecting RPC/RMI performance
 - Marshalling, unmarshalling, operations and despatch at the server
 - Data copying from application to kernel space to communication buffers.
 - Thread scheduling, context switching, and kernel entry.
 - Protocol processing for each protocol layer
 - Network access delays of connection setup, and network latency
- Concurrent and asynchronous invocations
 - Middleware or applications are made asynchronous if they don't block, waiting for reply to each invocation.

Communication and Invocation

- Most middleware such as CORBA, Java RMI, HTTP, is implemented over TCP.
 - Supporting universal availability, unlimited message size and reliable transfer.
 - Sun RPC (used in NFS) is implemented over both UDP and TCP and generally works faster over UDP.
- Research-based systems have implemented much more efficient invocation.
 - Firefly RPC (www.cdk3.net/oss)
 - Amoeba's `doOperation`, `getRequest`, `sendReply` primitives (www.cdk3.net/oss)
 - LRPC [Bershad et. al. 1990]

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

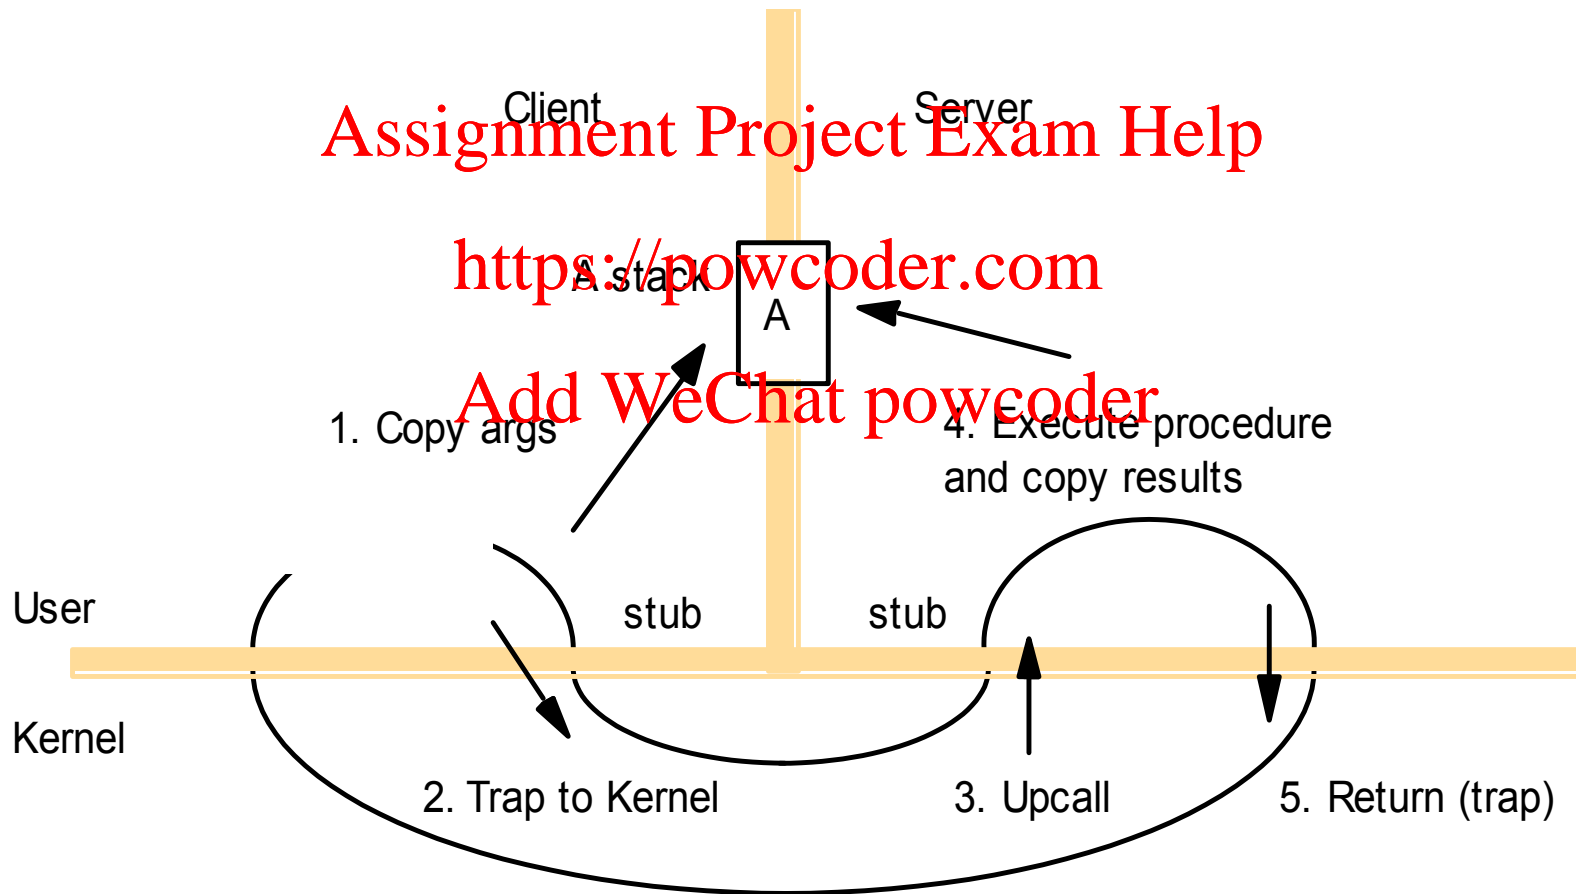
Communication and Invocation

□ LRPC (Lightweight RPC)

- Uses shared memory for interprocess communication, while maintaining protection of the two processes.
- Arguments copied only once rather than four times for conventional RPC.
 - Client stub stack to a message
 - The message to a kernel buffer
 - The kernel buffer to a server message
 - The message to the server stub stack
- Client threads can execute server code via protected entry points only.
- Up to 3 x faster for local invocations.

Communication and Invocation

□ LRPC (Lightweight RPC)



Communication and Invocation

□ Asynchronous invocations

- An asynchronous invocation is the one that is performed asynchronously with respect to the caller.
- An asynchronous invocation returns as soon as the invocation request message has been created and is ready for dispatch.
- Asynchronous invocation is applicable
 - When the client does not require response.
 - When the client uses a separate call to collect results.
- Middleware or applications are made asynchronous if they don't block, waiting for reply to each invocation.

Summary

- The combination of middleware and NOSs is a solution of autonomy and network-transparency.
- Modern OSs provide fundamental services for resource management, such as memory, process and thread, and communication management, on which middleware is built in support of distributed applications.
- Multiple threading is a cheaper way to implement concurrency than multiple processes, where Java is powerful tool for concurrent programming.
- Performance issues should be taken into account when designing/using middleware.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder